

DATASHEET



NVIDIA

NVIDIA DGX A100

The universal system for AI infrastructure.





NVIDIA



NVIDIA DGX-1

THE ESSENTIAL INSTRUMENT FOR AI RESEARCH

SYSTEM SPECIFICATIONS

NVIDIA DGX A100 320GB		
GPUs	8x NVIDIA A100 40GB Tensor Core GPUs	
GPU Memory	320 GB total	
Performance	5 petaFLOPS AI 10 petaOPS INT8	
NVIDIA NVSwitches	6	
System Power Usage	6.5 kW max	
CPU	Dual AMD Rome 7742, 128 cores total, 2.25 GHz (base), 3.4 GHz (max boost)	
System Memory	1TB	
Networking	Up to 8x Single-Port NVIDIA ConnectX-7 200 Gb/s InfiniBand 1 Dual-Port ConnectX-7 VPI 10/25/50/100/200 Gb/s Ethernet	Up to 8x Single-Port NVIDIA ConnectX-6 VPI 200 Gb/s InfiniBand 1 Dual-Port ConnectX-6 VPI 10/25/50/100/200 Gb/s Ethernet
Storage	OS: 2x 1.92TB M.2 NVME drives Internal Storage: 15TB (4x 3.84 TB) U.2 NVMe drives	
Software	DGX OS / Ubuntu / Red Hat Enterprise Linux / Rocky – Operating System NVIDIA Base Command – Orchestration, scheduling, and cluster management NVIDIA AI Enterprise – Optimized AI software	
Support	Comes with 3-year business-standard hardware and software support	
System Weight	271.5 lbs (123.16 kgs) max	
Packaged System Weight	359.7 lbs (163.16 kgs) max	
System Dimensions	Height: 10.4 in (264.0 mm) Width: 19.0 in (482.3 mm) max Length: 35.3 in (897.1 mm) max	
Operating Temperature Range	5–30 °C (41–86 °F)	

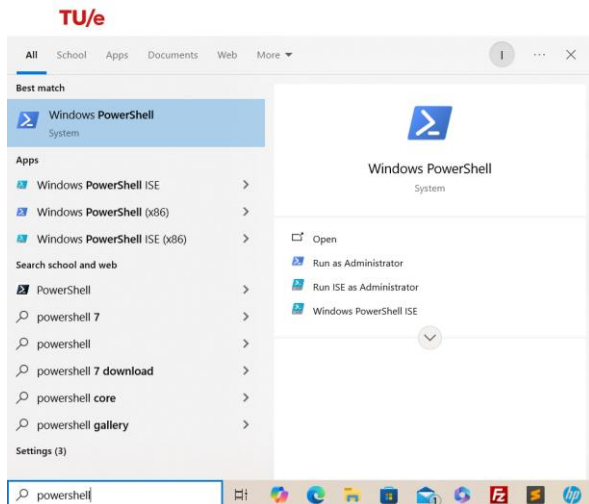


SYSTEM SPECIFICATIONS

GPUs	8X NVIDIA® Tesla® V100
Performance (Mixed Precision)	1 petaFLOPS
GPU Memory	256 GB total system
CPU	Dual 20-Core Intel Xeon E5-2698 v4 2.2 GHz
NVIDIA CUDA Cores	40,960
NVIDIA Tensor Cores (on Tesla V100 based systems)	5,120
Power Requirements	3,500 W
System Memory	512 GB 2,133 MHz DDR4 RDIMM
Storage	4X 1.92 TB SSD RAID 0
Network	Dual 10 GbE, 4 IB EDR
Operating System	Canonical Ubuntu, Red Hat Enterprise Linux
System Weight	134 lbs
System Dimensions	866 D x 444 W x 131 H (mm)
Packing Dimensions	1,180 D x 730 W x 284 H (mm)
Operating Temperature Range	5–35 °C

1. Accessing server

Pengguna windows: bisa dengan menggunakan **Windows Powershell**. Harap koneksi via VPN-BRIN terlebih dahulu apabila mengakses server dari jaringan eksternal BRIN.

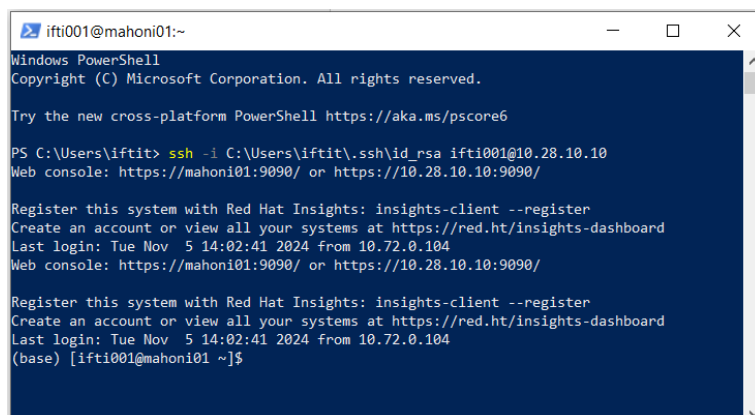


`<id_rsa location path>` lokasi di mana `id_rsa` disimpan di direktori lokal. Misalnya:
`C:\Users\<windows_username>\.ssh\id_rsa`

`<userid>` user intra BRIN (untuk user pengguna internal BRIN). Misalnya: `ifti001`

```
ssh -i <id_rsa location path> <userid>@10.28.10.10
```

Tampilan terminal setelah berhasil login ssh ke server.



Pengguna linux: bisa dengan Linux terminal.

2. Installation: Python libraries via anaconda and pip

Dokumentasi ini hanya mencakup cara instalasi library Python untuk riset di Machine Learning/Artificial Intelligence/Natural Language Processing. Apabila ada perbedaan cara maupun tipe modul library yang diinstal, harap mencari referensi tambahan yang tidak dibahas pada source repository terkait (misalnya, github).

Instalasi python libraries dilakukan setelah berhasil **login ssh di mahoni01**.

Catatan: Contoh penginstalan di sini adalah apabila user pengguna ingin menginstall library utama tanpa dependen ke repository github tertentu. Sebaliknya, harap menggunakan referensi dari github atau repository acuan.

- Instalasi miniconda3 untuk penginstalan via conda

```
mkdir -p ~/miniconda3

wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh
-O ~/miniconda3/miniconda.sh

bash ~/miniconda3/miniconda.sh -b -u -p ~/miniconda3
rm ~/miniconda3/miniconda.sh
source ~/miniconda3/bin/activate
conda init
```

- Instalasi library via conda environment, referensi pytorch <https://pytorch.org/get-started/locally/>

```
conda create --name my_env python=3.11
conda activate my_env

conda install pytorch torchvision torchaudio pytorch-cuda=12.1 -c pytorch -c
nvidia

pip install transformers evaluate pandas matplotlib wandb scikit-learn
```

3. Writing job script

- Contoh bash script. Simpan dengan nama misal: **run_script.sh**

```
#!/bin/bash
source /home/ifti001/miniconda3/etc/profile.d/conda.sh
conda activate my_env

python python_script1.py
```

- Contoh python script yang dibaca oleh bash script. Simpan dengan nama **python_script.py**

```
import os
import sys
import numpy as np
import pandas as pd

# Please use transformers==4.45.2
import transformers
import torch

def main():

    # loading model for Huggingface hub
    model_id = "aisingapore/gemma2-9b-cpt-sea-lionv3-instruct"

    pipeline = transformers.pipeline(
        "text-generation",
        model=model_id,
        model_kwargs={"torch_dtype": torch.bfloat16},
        device_map="auto",
    )

    messages = [
        {"role": "user", "content": "Apa sentimen dari kalimat berikut ini?\nKalimat: Buku ini sangat membosankan.\nJawaban: "},
    ]

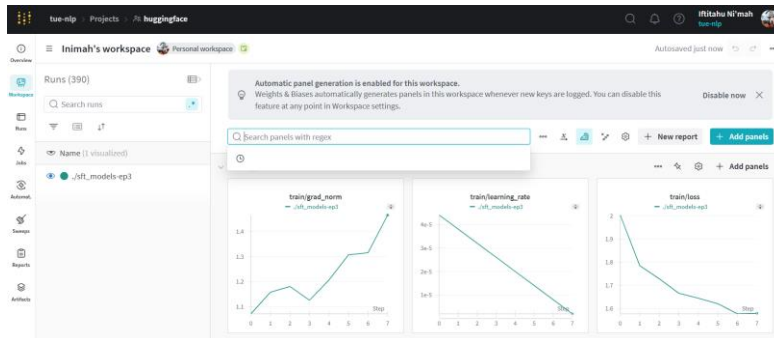
    outputs = pipeline(
        messages,
        max_new_tokens=256,
    )
    print(outputs[0]["generated_text"][-1])

if __name__ == '__main__':
    main()
```

Q: Bagaimana apabila saya ingin melihat progress training model secara interaktif?

A: Huggingface sudah mengintegrasikan module **Trainer** dengan <https://wandb.ai/>

Harap mengacu pada referensi yang ada. Contoh tampilan wandb untuk progress training model:



4. Transferring file

Dari server ke lokal direktori.

Harap menjalankan script di **Windows Powershell** untuk pengguna windows (tidak perlu login ssh ke server). Contoh:

```
scp -r ifti001@10.28.10.10:/home/ifti001/<FILEPATH> ./<LocalDIR>
```

Dari direktori lokal ke server

```
scp -r .\<LOCAL-FILEPATH> ifti001@10.28.10.10:/home/ifti001/
```

5. Submit job script

```
sbatch --partition=short --gres=gpu:1 --mem=4G --odelist=a100  
<BASH_SCRIPT.sh>
```

--partition ['short', 'medium'] Ekspektasi timeline atau lama job yang akan disubmit: 'short' = 24 jam; 'medium' = 3 hari.

--gres=gpu:1 Apakah job yang disubmit akan menggunakan device 'gpu' atau 'cpu'. Untuk preprocessing atau pembacaan data dari dan ke disk space masing-masing user, utilisasinya menggunakan 'cpu'. Untuk proses inference model yang tidak memerlukan loading model di 'gpu', utilisasinya bisa diarahkan ke 'cpu'.

--odelist [a100, a1] Node partisi yang akan digunakan.

<BASH_SCRIPT.sh> nama job script yang akan disubmit dalam antrian slurm di server.

6. Job Audit dan Perintah Dasar Linux

- Melihat antrian job

```
squeue -a
```

- Meng-cancel job

```
scancel <JOBID>
```

- Melihat status GPU yang tersedia

```
sinfo
```

- Melihat lokasi direktori

```
pwd
```

- Copy, move file ke directory lain

```
cp -rf <CURRENT-FILEPATH> <TARGET-DIR>  
mv <CURRENT-FILEPATH> <TARGET-DIR>
```

- Remove file

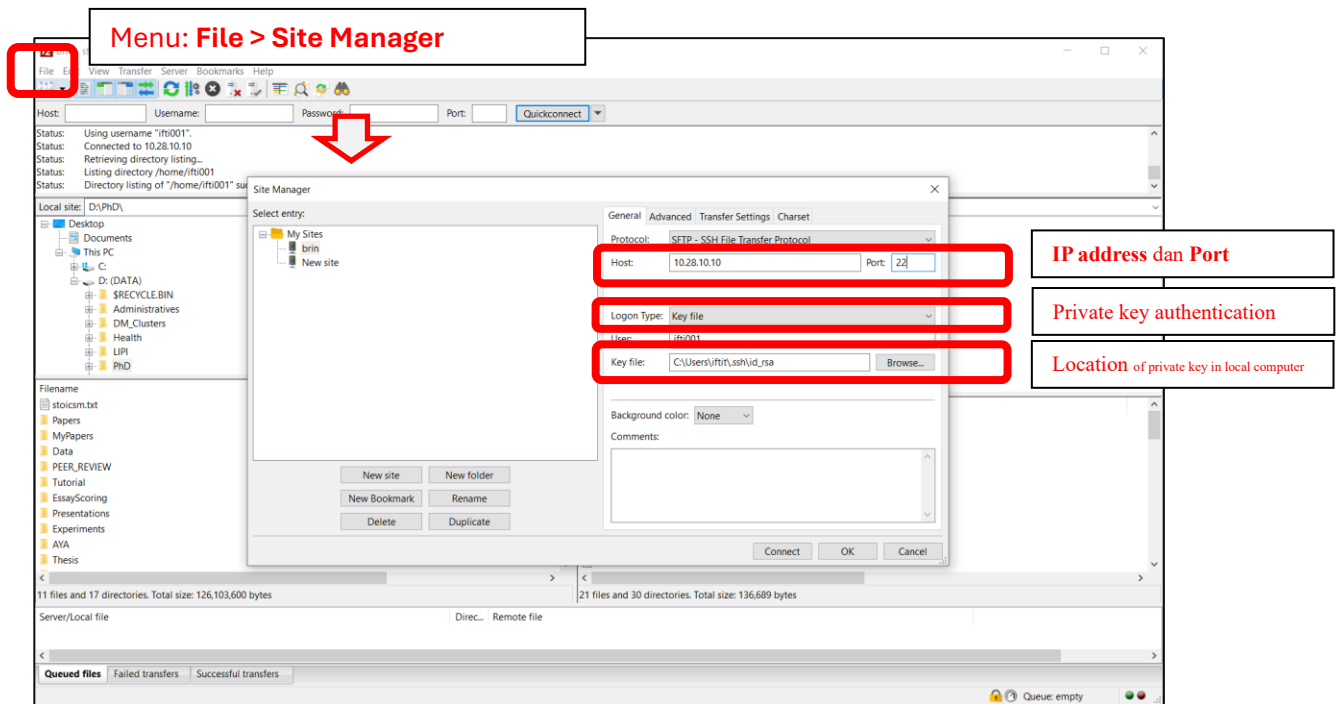
```
rm -rf <CURRENT-FILEPATH>
```

7. Setting Filezilla dengan private RSA key

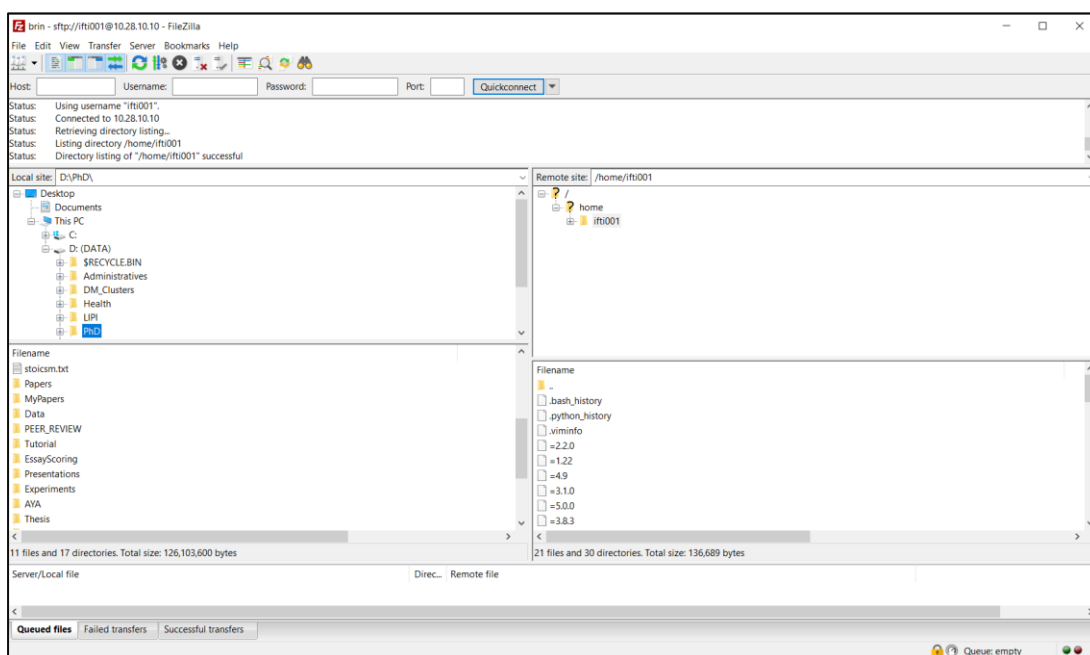
Download dan install Filezilla client for windows:

<https://filezilla-project.org/download.php?platform=win64>

Konfigurasi **File > Site Manager**



Tampilan ketika berhasil connect ke server HPC via Filezilla. Setelah ini bisa transfer file dengan drag-and-drop atau klik kanan dari dan ke server.



8. Model save checkpoints and continue training

Catatan: Silahkan disesuaikan dengan domain riset masing-masing menurut referensi yang relevant.

- Dengan Pytorch save and load (<https://wandb.ai/wandb/common-ml-errors/reports/How-to-save-and-load-models-in-PyTorch--VmlldzozMjg0MTE> ; https://pytorch.org/tutorials/beginner/saving_loading_models.html)

▼ Save and load your PyTorch model from a checkpoint

In most machine learning pipelines, saving model checkpoints periodically or based on certain conditions is essential. This practice allows you to resume training from the latest or best checkpoint, ensuring continuity in case of interruptions. Checkpoints are also useful for fine-tuning and evaluating model performance at different stages.

When saving a checkpoint, saving only the model's `state_dict` is not sufficient. You should also save the optimizer's `state_dict`, the last epoch number, the current loss, and any other relevant information needed to seamlessly resume training.

▼ Save a PyTorch model checkpoint

```
torch.save({'epoch': EPOCH,
           'model_state_dict': model.state_dict(),
           'optimizer_state_dict': optimizer.state_dict(),
           'loss': LOSS},
          'save/to/path/model.pth')
```

▼ Load a PyTorch model checkpoint

```
model = MyModelDefinition(args)
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

checkpoint = torch.load('load/from/path/model.pth')
model.load_state_dict(checkpoint['model_state_dict'])
optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
epoch = checkpoint['epoch']
loss = checkpoint['loss']
```

▼ Gotchas:

- When resuming training, remember to call `model.train()` to ensure the model is set to training mode.
- For inference after loading a checkpoint, call `model.eval()` to set the model to evaluation mode.

- Dengan accelerator: https://huggingface.co/docs/accelerate/en/usage_guides/checkpoint
- Dengan PEFT adapter continue training: <https://colab.research.google.com/drive/12pMorxvLV-VwjuNBM76L4xXnzVYg57iB?usp=sharing>